

Covariant programming : capturing the variation between SIMD and SIMT in a single code

WAS

```
LatticeFermionView a,b,c;  
accelerator_for(ss,volume, {  
    a[ss] = b[ss] + c[ss] ;  
});
```

NOW

```
LatticeFermionView a,b,c;  
accelerator_for(ss,volume,Spinor::Nsimd(), {  
    coalescedWrite(a[ss], b(ss) + c(ss) );  
});
```

On GPU accelerator for sets up a volume \times Nsimd thread grid.

- Each thread is responsible for one SIMD lane

On CPU accelerator for sets up a volume OpenMP loop.

- Each thread is responsible for Nsimd() SIMD lanes

Per-thread datatypes inside these loops cannot be hardwired.

C++ auto and decltype use the return type of operator () to work out computation variables in architecture dependent way.

**Single source high performance kernels are optimal
on BOTH CPU and CUDA**

